



Here are the pre-processing graph results.


Heatmap of Attribute Correlation




AT




V




AP



RH





PE



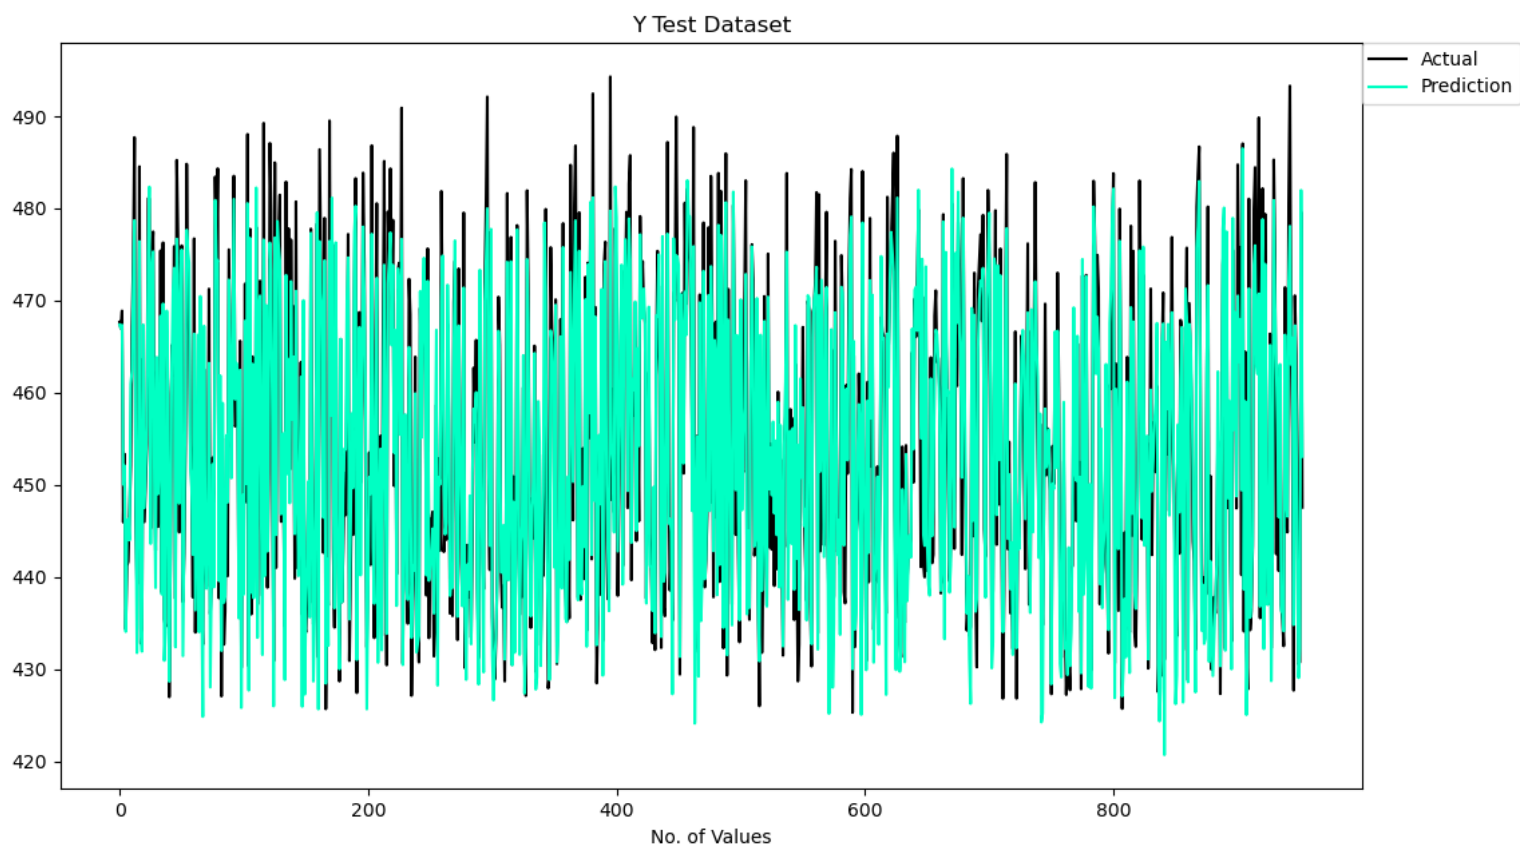
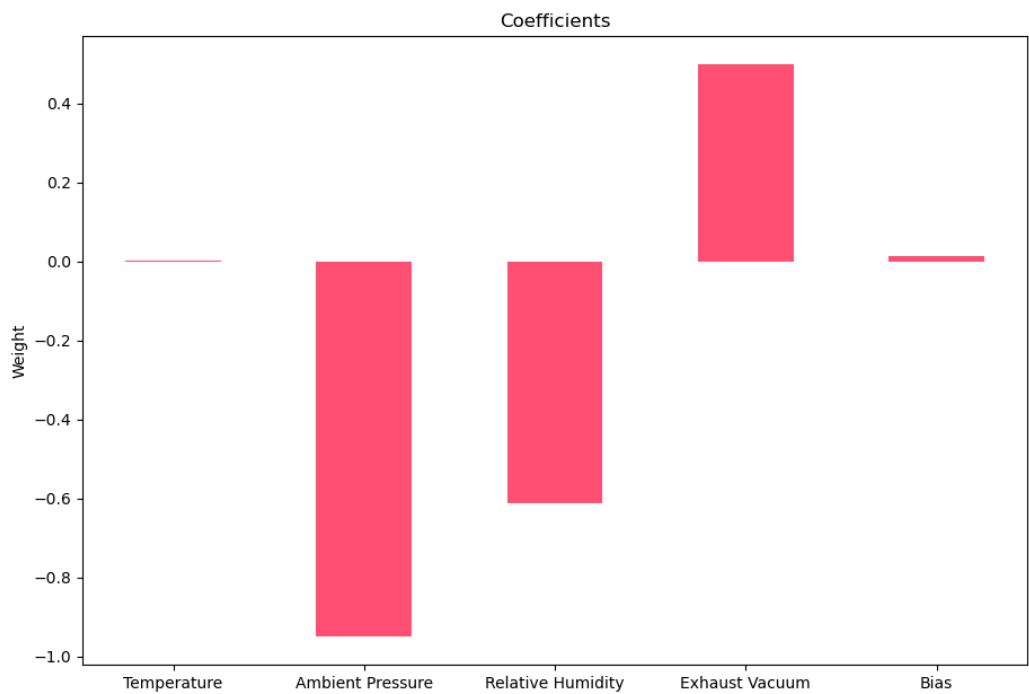
Plots of the attributes. The last plot is of the output value.

Part1 Plots for the trial with the best parameters found. State = 5.





Part2 Plots for the trial with the best parameters found. State = 5.



Part 1 of Logs

Tuning the Enhanced Gradient Descent Model for Best Parameters

Part 1 of Enhanced Gradient Descent Parameters Used: State: 5 Alpha: .001 Beta1: .9 Beta2: .999 Epsilon: 10^-8 m = 0 v = 0 Iterations: 5000 Train Split: 80.0 % Test Split: 20.0 % Coefficients: [0.4374861331222539, -0.6123631130942228, -0.26003467414341763, 0.44559254221717787, 0.3878470184438752] Train Accuracy: Mean Squared Error: 81.5735800926913 R^2 Value: 0.48285457574898716 Test Accuracy: Mean Squared Error: 80.98996701556176 R^2 Value: 0.49817895252641486
Part 1 of Enhanced Gradient Descent Parameters Used: State: 5 Alpha: .001 Beta1: .9 Beta2: .999 Epsilon: 10^-8 m = 0 v = 0 Iterations: 8000 Train Split: 80.0 % Test Split: 20.0 % Coefficients: [0.48719170002575046, -1.118172756881196, -0.5029630792192917, 0.4937123219603581, 0.03896292313804325] Train Accuracy: Mean Squared Error: 29.51493201645292 R^2 Value: 0.8844960353362512 Test Accuracy: Mean Squared Error: 29.646878731320733 R^2 Value: 0.8864923238960015
Part 1 of Enhanced Gradient Descent Best one found. Parameters Used: State: 5 Alpha: .001 Beta1: .9 Beta2: .999 Epsilon: 10^-8 m = 0 v = 0 Iterations: 10000 Train Split: 80.0 % Test Split: 20.0 % Coefficients: [0.5144726250635316, -1.423992626029842, -0.39589981822794335, 0.5009286790132073, -0.058344123852125625] Train Accuracy: Mean Squared Error: 26.22033816483604 R^2 Value: 0.9019748858552826 Test Accuracy: Mean Squared Error: 26.42235572271788 R^2 Value: 0.9030012797938725
Part 1 of Enhanced Gradient Descent Drops attribute: 'AT' (this can be verified with the heatmap showing it's correlation with 'V') Parameters Used: State: 5 Alpha: .001 Beta1: .9 Beta2: .999 Epsilon: 10^-8 m = 0 v = 0

Iterations: 10000

Train Split: 80.0 % | Test Split: 20.0 %

Coefficients:

[0.4648212353511957, -1.0136507011995473, 0.4906395683339943, 0.16063057719700888]

Train Accuracy:

Mean Squared Error: 57.38900595212531

R² Value: 0.7500325370547691

Test Accuracy:

Mean Squared Error: 55.812871327546695

R² Value: 0.7628366986014008

Part 1 of Enhanced Gradient Descent

Parameters Used:

State: 5

Alpha: .001 | Beta1: .9 | Beta2: .999 | Epsilon: 10⁻⁸ | m = 0 | v = 0

Iterations: 10000

Train Split: 90.0 % | Test Split: 10.0 %

Coefficients:

[0.5140621380109054, -1.422759204876915, -0.39610464580669436, 0.5009524502521513, -0.05821683974215756]

Train Accuracy:

Mean Squared Error: 26.20013185328805

R² Value: 0.9024861152060023

Test Accuracy:

Mean Squared Error: 26.86479385463498

R² Value: 0.8985095209493589

Note:

- A key thing to mention here is that the Adam Optimizer is recommended to have the same variable values for Alpha, Beta1, Beta2, Epsilon, initial m, and initial v. These standard values are used through most implementations of using the Adam optimization. After trying to tune the parameters for the Enhanced Gradient Descent Model here were the best parameters found: iterations = 10000, train/test split = 90/10, and keeping all attributes.

Part 2 of Logs

Comparing Gradient Descent & Sklearn Stochastic Gradient Descent Model

Part 1 of Enhanced Gradient Descent

Parameters Used:

State: 1

Alpha: .001 | Beta1: .9 | Beta2: .999 | Epsilon: 10^{-8} | m = 0 | v = 0

Iterations: 10000

Train Split: 80.0 % | Test Split: 20.0 %

Coefficients:

[0.5135360955599416, -1.4129702682082277, -0.4030913398665183, 0.5012177628402052, -0.0601577196111157]

Train Accuracy:

Mean Squared Error: 25.907627020055898

R² Value: 0.9030119875133472

Test Accuracy:

Mean Squared Error: 28.014235643608277

R² Value: 0.8973077744212778

Part 2 of Enhanced Gradient Descent

Parameters Used:

State: 1

Iterations: 10000

Learning Rate: 1e-06

Coefficients:

[6.33944060e-04 -1.04579006e+00 -5.69972842e-01 5.00032224e-01 -1.42396194e-03]

Train Accuracy:

Mean Squared Error: 29.96784033730602

R² Value: 0.8827791382973325

Test Accuracy:

Mean Squared Error: 35.768615316443594

R² Test: 0.8706809198549961

Part 1 of Enhanced Gradient Descent

Parameters Used:

State: 2

Alpha: .001 | Beta1: .9 | Beta2: .999 | Epsilon: 10^{-8} | m = 0 | v = 0

Iterations: 10000

Train Split: 80.0 % | Test Split: 20.0 %

Coefficients:

[0.5146539289431077, -1.4216915288677878, -0.3944771045499932, 0.5007776345920979, -0.05787278024584849]

Train Accuracy:

Mean Squared Error: 26.35812790742656

R² Value: 0.9008578375693331

Test Accuracy:

Mean Squared Error: 25.894430060307297

R² Value: 0.9052769558060673

Part 2 of Enhanced Gradient Descent

Parameters Used:

State: 2

Iterations: 10000

Learning Rate: 1e-06

Coefficients:

[8.26639238e-04 -1.17934925e+00 -5.03116141e-01 5.00811753e-01 -2.21790217e-02]

Train Accuracy:

Mean Squared Error: 29.329251229280484

R² Value: 0.8865946508768744

Test Accuracy:

Mean Squared Error: 28.93303057712341

R² Test: 0.8899080573047872

Part 1 of Enhanced Gradient Descent

Parameters Used:

State: 3

Alpha: .001 | Beta1: .9 | Beta2: .999 | Epsilon: 10⁻⁸ | m = 0 | v = 0

Iterations: 10000

Train Split: 80.0 % | Test Split: 20.0 %

Coefficients:

[0.514850161615456, -1.4175520120905034, -0.4032907854138991, 0.5010969321402667, -0.05667334052436697]

Train Accuracy:

Mean Squared Error: 26.049164151569435

R² Value: 0.90280841347468

Test Accuracy:

Mean Squared Error: 27.403502699222685

R² Value: 0.9015783901380442

Part 2 of Enhanced Gradient Descent

Parameters Used:

State: 3

Iterations: 10000

Learning Rate: 1e-06

Coefficients:

[5.79192509e-04 -9.85495454e-01 -5.98517524e-01 4.98454730e-01 7.34514243e-03]

Train Accuracy:

Mean Squared Error: 31.1965394854964

R² Value: 0.8778546902838902

Test Accuracy:

Mean Squared Error: 32.31482406991121

R² Test: 0.8783472782239139

Part 1 of Enhanced Gradient Descent

Parameters Used:

State: 4

Alpha: .001 | Beta1: .9 | Beta2: .999 | Epsilon: 10⁻⁸ | m = 0 | v = 0

Iterations: 10000

Train Split: 80.0 % | Test Split: 20.0 %

Coefficients:

[0.5143122919765742, -1.424814866085679, -0.3971354213196689, 0.5011520085825308, -0.060557465214266824]

Train Accuracy:

Mean Squared Error: 26.236927421839752

R² Value: 0.9020796768035431

Test Accuracy:

Mean Squared Error: 26.35186963702148

R² Value: 0.9030362793154947

Part 2 of Enhanced Gradient Descent

Parameters Used:

State: 4

Iterations: 10000

Learning Rate: 1e-06

Coefficients: [6.18053193e-04 -1.02958395e+00 -5.75053426e-01 4.99599192e-01 -1.20913344e-03] Train Accuracy: Mean Squared Error: 30.428115367405606 R^2 Value: 0.8811448017908418 Test Accuracy: Mean Squared Error: 29.810869047270845 R^2 Test: 0.8834203548265686
Part 1 of Enhanced Gradient Descent Parameters Used: State: 5 Alpha: .001 Beta1: .9 Beta2: .999 Epsilon: 10^-8 m = 0 v = 0 Iterations: 10000 Train Split: 80.0 % Test Split: 20.0 % Coefficients: [0.5144726250635316, -1.423992626029842, -0.39589981822794335, 0.5009286790132073, -0.058344123852125625] Train Accuracy: Mean Squared Error: 26.22033816483604 R^2 Value: 0.9019748858552826 Test Accuracy: Mean Squared Error: 26.42235572271788 R^2 Value: 0.9030012797938725
Part 2 of Enhanced Gradient Descent Parameters Used: State: 5 Iterations: 10000 Learning Rate: 1e-06 Coefficients: [5.20305073e-04 -9.22055912e-01 -6.25737931e-01 4.97391759e-01 1.76156170e-02] Train Accuracy: Mean Squared Error: 33.3083715750111 R^2 Value: 0.8684416548402241 Test Accuracy: Mean Squared Error: 33.33351926639194 R^2 Test: 0.8712252043979538

Notes:

- The chosen parameters for the SGD-Regressor kept the iterations the same as the ones used for the Enhanced Gradient Descent, set the learning rate to 1e-06 since it was the only Learning Rate that actually was optimal, set the loss function to “squared_loss”, and did not set any penalty.
- The states here are kept the same to properly compare the models.
- The obvious conclusion to be made here is that from the chosen parameters for both models, the Enhanced Gradient Descent model that implemented the Adam optimizer performed better than the sklearn SGD-Regressor model. Lower Mean Squared Error value and a higher R^2 value.

Answers to Questions

Are you satisfied that you have found the best solution? Explain.

- Overall, we are satisfied with the results we obtained from our enhanced version of vanilla Gradient Descent. The reason for this is that the accuracy achieved is something that can be interpreted as a success. Of course, there are other optimizations that can be used for the model to achieve even better results such as Normalizing or Standardizing the data before sending it through the model. There could also be other pre-processing analysis that can help modify the data to be more optimal as well. Since the accuracy seems sufficient and it is not only comparable to the SGD-Regressor model but achieves better results, we can be satisfied that we found a good solution.

Are you satisfied that the package has found the best solution? Explain.

- We feel that there is a lot more that could be done with the SGD-Regressor package that we didn't implement. There were a lot more parameters and other resources that could have helped fine tune the model. On top of Normalizing or Standardizing the data, we could have spent more time understanding all the inputs into the model that were offered. Furthermore, we are trying to compare the Adam optimizer to the SGD-Regressor model so it would be difficult to alter parameters without just having two almost incomparable models. Due to this concern we kept the SGD-Regressor model as simple as possible in order to have an easier understanding of what was directly being compared.